Towards Automated Choreographing of Web Services using Planning

Guobing Zou[†‡§]**, Yixin Chen**[‡]**, You Xu**[‡]**, Ruoyun Huang**[‡]**, Yang Xiang**[§]
[†]School of Computer Engineering and Science, Shanghai University, China
[‡]Dept. of Computer Science and Engineering, Washington University in St. Louis, USA
[§]Dept. of Computer Science and Technology, Tongji University, China
guobingzou@gmail.com, {chen, yx2, rh11}@cse.wustl.edu, shxiangyang@tongji.edu.cn

Abstract

For Web service composition, choreography has recently received great attention and demonstrated a few key advantages over orchestration such as distributed control, fairness, data efficiency, and scalability. Automated design of choreography plans, especially distributed plans for multiple roles, is more complex and has not been studied before. Existing work requires manual generation assisted by model checking. In this paper, we propose a novel planning-based approach that can automatically convert a given composition task to a distributed choreography specification. Although planning has been used for orchestration, it is difficult to use planning for choreography, as it involves decentralized control, concurrent workflows, and contingency. We propose a few novel techniques, including compilation of contingencies, dependency graph analysis, and communication control, to handle these characteristics using planning. We theoretically show the correctness of this approach and empirically evaluate its practicability.

Introduction

Orchestration and choreography are two ways for Web service composition. A WSC problem can be described from the view of a single participant using orchestration or from a global perspective using choreography (Barker, Walton, and Robertson 2009). Service orchestration refers to an executable business process that has a central controller process to coordinate all of the participating Web services (Peltz 2003). On the contrary, service choreography does not have an orchestrator, but all of the participating Web services collaborate with each other in order to achieve a shared goal. The Web Services Choreography Description Language (WS-CDL) is an XML-based W3C candidate language for describing collaborations of service choreography.

It has been widely recognized that choreography has certain key advantages over orchestration, including less data transfer, robustness under server failure, fairness, and avoidance of deadlocks. Although service choreography is widely advocated, automated generation of service choreography specification has not been studied. WS-CDL describes choreography from a global view in a single master plan.

It is realized that such a global view is not sufficient and choreography plans should be distributed (Qiu et al. 2007). Although a global view such as defined by WS-CDL is helpful, at the execution end, each role should have a "local plan" such as the MAP protocol (Barker, Walton, and Robertson 2009) that specifies what it needs to do from its individual perspective. Designers ensure that the collaboration of the multiple roles correctly achieves the global goals.

Some recent research focuses on new languages for distributed plans in choreography, but relies on users to manually write the specifications. Examples include the Multiagent Protocols (MAP) (Barker, Walton, and Robertson 2009), and WS-CDL+ (Kang, Wang, and Hung 2007). They do not solve the problem of how to generate these specifications, although there are works that verify the protocol using model checking (Barker, Walton, and Robertson 2009; Yang et al. 2008). Due to the complexity of decentralized logics, manually completing such a specification for service choreography can be tedious and error prone. Although model checking can be used to verify the plan, designing by trial-and-error is labor intensive.

In this paper, we propose a novel approach that can automatically generates a distributed choreography plan based on automated planning. Although planning has been used for orchestration, it is difficult for choreography for several challenges. While planning is suitable for constructing the composition plan from the view of a single party in orchestration, choreography needs distributed and coordinated plans for multiple participants, making planning more difficult. Also, choreography needs to support asynchronous communication between peers and contingent plans that depend on the outcomes of Web services.

We propose a novel planning-based framework to address the above challenges. Our framework compiles a repository of related Web services, along with user-defined contingencies, into a planning domain in Planning Domain Definition Language (PDDL). This process compiles away the contingency on action outcomes into a deterministic planning problem. A planner is used to find a master plan which provides a global view of the choreography. We further propose a decentralization scheme which uses a dependency analysis to support the synthesis of local plans, one for each peer. By addressing three main features in choreography, including *choices*, *parallelism*, and *communication*, our approach
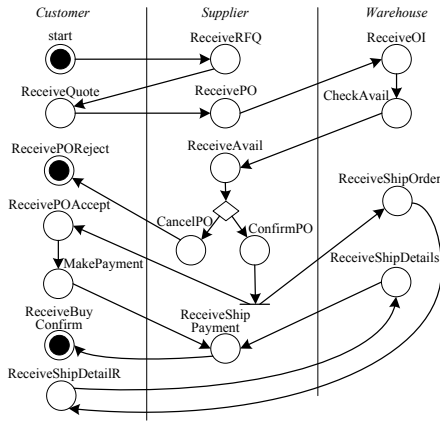
Figure 1: The interaction flow of the three roles.

is correct in the sense that all possible distributed execution sequences carried out by the participating peers are valid under the centralized master plan.

We implemented a prototype system and conducted extensive experiments on 81,464 WSDL services in 18 groups of large scale service repositories. The experimental results validated the feasibility and efficiency of our proposed work.

A Running Example

Our example has three Web services: customer, supplier and warehouse. Figure 1 shows an ordering process.

The **Customer** sends a quote request to the **Supplier** on a given product, and then the **Supplier** receives the request (ReceiveRFQ) and replies. After receiving it (ReceiveQuote), the **Customer** sends a request to the **Supplier**. Once receiving an order (ReceivePO), the **Supplier** sends order information to the **Warehouse**, who receives it (ReceiveOI), checks its availability (CheckAvail), and replies to the **Supplier**. When the **Supplier** receives it (ReceiveAvail) from the **Warehouse**, it makes a decision for the order request.

There are two possibilities. If a product is unavailable, the **Supplier** cancels it (CancelPO) and notifies the **Customer**. Otherwise, the **Supplier** accepts it (ConfirmPO). After receiving the acceptance (ReceivePOAccept), the **Customer** makes a payment (MakePayment) and replies a payment confirmation to the **Supplier**. Meanwhile, after receiving a shipping order request (ReceiveShipOrder), the **Warehouse** sends a request to the **Customer** for shipping details. Once the **Customer** receives it (ReceiveShipDetailR), it replies shipping details to the **Warehouse**. After receiving a payment and shipping confirmation (ReceiveShipPayment), the **Supplier** replies to a purchase confirmation to **Customer**.

Problem Formulation

We first formulate our problem using a simplified model. In this paper, we focus on understanding the fundamental principles of automated choreography by planning, and simplify or omit certain issues in Web service composition, such as ontology (Agarwal et al. 2005) and background theory (Hoffmann, Bertoli, and Pistore 2007). The theory and algorithms can be readily incorporated with other techniques

and extended to more complex systems. In the experimental results section, we present an evaluation using real-world WSDL Web services and MAP for choreography description. As we will see, the projection from our theoretical model to these languages are straightforward.

Definition 1 (Web Service). *A Web service w consists of a set of operations, denoted as $w = \{op_1, op_2, \cdots\}$, where $\forall op \in w$ is a 2-tuple $\langle I, O \rangle$, where $I = \{I^1, I^2, \cdots\}$ is a set of input interface parameters. Similarly, $O = \{O^1, O^2, \cdots\}$ is a set of output interface parameters. We use $op.I$ and $op.O$ to denote I and O in op, respectively. For each interface parameter x, we use $Dom(x)$ to denote its possible values and $x.value \in Dom(x)$ to denote the value of x.*

Note that each Web service plays a role that can perform a set of operations. A service repository is a set of services.

Definition 2 (Contingency). *For an operation op with $op.I = \{I^1, I^2, \cdots\}$, a contingency is a tuple $c = (op, I^i, pre^i)$ where $pre^i \subset Dom(I^i)$.*

A contingency $c = (op, I^i, pre^i)$ means that to invoke op, we need its input parameter I^i to take the values in pre^i, instead of any value in $Dom(I^i)$.

We define a **service state** as a set of interface parameters $Q = \{x^1, x^2, \cdots\}$ that are defined and assigned values. We assume that parameters not in Q are unavailable at the state.

Definition 3 (Applicability). *Without contingency, an operation op is applicable at a service state Q if $op.I \subseteq Q$. We denote this as $Q \succ op$. An operation op under contingency $c = (op, I^i, pre^i)$ (denoted as $c \triangleright op$) is applicable at Q if $op.I \subseteq Q$ and $op.I^i.value \in pre^i$. We denote this as $Q \succ c \triangleright op$.*

When an applicable operation op or $c \triangleright op$ is applied to Q, the resulting state $Q' = Q \oplus op$ (or $Q' = Q \oplus c \triangleright op$) is $Q' = Q \cup op.O$, in which the values of the parameters in $op.O$ are set by executing op. An **execution sequence** is an ordered list $L = (o_1, \cdots, o_m)$ where each element is either an operation op or an operation with contingency $c \triangleright op$. Applying a sequence L to a service state Q results in $Q' = Q \oplus L = (\cdots ((Q \oplus o_1) \oplus o_2) \cdots \oplus o_m)$ if every step is applicable (otherwise $Q \oplus L$ is undefined).

Definition 4 (Service Choreography Problem (SCP)). *A SCP is defined by (W, C, r_{in}, r_{out}): 1) a service repository $W = \{w_1, \cdots, w_N\}$, 2) a set of contingencies $C = \{c_1, \cdots, c_{N_c}\}$, 3) an input parameter set $r_{in} = \{r^1, r^2, \cdots\}$, and 4) an output parameter set $r_{out} = \{q^1, q^2, \cdots\}$.*

Example 1. Following our running example, its SCP has $W = \{Customer, Supplier, Warehouse\}$. An example operation is $ReceiveRFQ = \{I, O\}$, where $I = \{pid, pid\_name\}$, $O = \{pid\_price\}$. Another operation is $ConfirmPO = \{I, O\}$, where $I = \{po\_avail\}$, $O = \{po\_accept, shiporderR\}$. In particular, $Dom(po\_avail) = \{avail, not\_avail\}$, while po_accept and $shiporderR$ represent an order acceptance and shipping order request, respectively.

The contingency set consists of $C = \{c_1, c_2\}$, where $c_1 = (ConfirmPO, po\_avail, \{avail\})$ and $c_2 =$

$(CancelPO, po\_avail, \{not\_avail\})$. The initial choreography inputs $r_{in} = \{pid, pid\_name, pid\_quantity\}$. The goal is $r_{out} = \{purchase\_confirm, po\_order\_reject\}$, which has a purchase confirmation $purchase\_confirm$ and an order rejection po_order_reject. ∎

Note that here we assume that the user can specify multiple possible goal states he is interested in. Typically, the users are knowledgeable of these multiple choreography goals, because they are advanced model developers of Web service composition (Ponnekanti and Fox 2002). A choreography should consider all of the contingencies and give service choreography plans that can handle the various goals. In our example, the user specifies both $purchase\_confirm$ and po_order_reject as goals, so that the choreographer can find plans contingent on the availability of the product.

Given a SCP=(W, C, r_{in}, r_{out}), a **choreography master plan** is any expression P defined by the following language.

$$
\begin{array}{llll}
P & ::= & op & (op \in w \in W) \\
& | & talk(i,j) & (role\ i\ talks\ to\ j) \\
& | & P; P & (sequential) \\
& | & P \parallel P & (parallel) \\
& | & c \triangleright P\ \textsf{or}\ P & (c \in C, choice)
\end{array}
$$

where op is an operation in a Web service, and c is a contingency. We allow only one contingency for each **or** choice to simplify the presentation, although it is easy to extend to multiple contingencies. Our language is similar to other choreography description (Barker, Walton, and Robertson 2009; Qiu et al. 2007), except that we explicitly introduce contingency in our definition. The $talk$ action is applicable to any service state and brings no change to the state.

Given a service state Q, let $t(P)$ denote all of the possible execution sequences from Q, we can define:

$$
\begin{array}{rcl}
t(op) & = & \{(op)\} \\
t(talk(i,j)) & = & \{(talk(i,j))\} \\
t(P_1; P_2) & = & t(P_1) \circ t(P_2) \\
t(P_1 \parallel P_2) & = & t(P_1) \bowtie t(P_2) \\
t(c \triangleright P_1\ \textsf{or}\ P_2) & = & \left\{ \begin{array}{ll} t(P_1), & \text{if } Q \succ c \triangleright P_1; \\ t(P_2), & \text{otherwise.} \end{array} \right.
\end{array}
$$

where $\circ$ denotes the concatenation of two trace sets (i.e. $A \circ B = \{(a,b)|a \in A, b \in B\}$), and $\bowtie$ is the interleaving of two trace sets (the definition is standard and omitted); $Q \succ c \triangleright P_1$ should be understood as $Q \succ c \triangleright op$ for any op that can be the first operation in a sequence in $t(P_1)$.

Definition 5 (Centralized Solution). *A centralized solution to a SCP=(W, C, r_{in}, r_{out}) is a choreography master plan P such that for every sequence $L \in t(P)$, $r_{in} \oplus L$ is defined and $r_{in} \oplus L \supseteq r_{out}$.*

In SCP, there are multiple roles, each corresponds to a $w \in W$. A **local plan** is any expression R defined by the following language.

$$
\begin{array}{llll}
R & ::= & op & (op \in w_i) \\
& | & send(ch, i, j) & (send\ to\ role\ j) \\
& | & recv(ch, j, i) & (receive\ from\ role\ j) \\
& | & R; R & (sequential) \\
& | & R \parallel R & (parallel) \\
& | & c \triangleright R\ \textsf{or}\ R & (c \in C^i, choice)
\end{array}
$$

where $C^i \subseteq C$ is the set of contingencies related to w_i (i.e. C^i includes those c whose operation op is in w_i); and ch is a unique channel ID for each $send/recv$ pair. Like $talk$, $send$ and $recv$ are applicable to any service state.

Given a service state Q, for role w_i with a local plan R_i, we can similarly define $t(R_i)$, the set of possible execution sequences of w_i from Q. A **distributed choreography plan** R is a set of local plans R_i, one for each role $w_i \in W$. Then, we can define the set of **combination sequences** as

$$\mathbf{C}(R) = \{\bowtie^* (L_1, \cdots, L_N)|L_i \in t(R_i), i = 1..N\},$$

where $\bowtie^*$ denotes any interleaving of N sequences subject to one constraint: $send(ch, i, j)$ is always before $recv(ch, j, i)$ for any ch, i and j.

Definition 6 (Distributed Solution). *A distributed solution to a SCP=(W, C, r_{in}, r_{out}) is a distributed choreography plan R such that for every sequence $L \in \mathbf{C}(R)$, $r_{in} \oplus L$ is defined and $r_{in} \oplus L \supseteq r_{out}$.*

Definition 7 (Equivalence). *A centralized solution is equivalent to a distributed solution when their sequence sets contain identical sequences, ignoring $talk$, $send$, and $recv$.*

A closely related equivalence has been studied in (Qiu et al. 2007). We comment that WS-CDL can be viewed as an extended language for master plans, while MAP (Barker, Walton, and Robertson 2009) and the Role Language in (Qiu et al. 2007) are examples of languages for distributed plans.

Automated Choreography by Planning

We develop an approach that can correctly generate a distributed solution for a SCP. It has a few major steps. 1) Translate an SCP into a PDDL planning problem, which complies away action contingency. 2) Solve the planning problem using an automated planner to obtain a solution plan. 3) Perform a dependency analysis on the plan to build a dependency graph (DP) and mark DP to generate a master plan P. 4) Project P to a distributed plan R based on the DP. We will also show that R is a distributed solution to the SCP, by showing that P is a centralized solution and that R is equivalent to P.

Planning formulation of SCP

A classical planning problem π is a tuple $\pi = (A, F, \mathcal{I}, G)$, where A is the set of actions, F the set of facts, $\mathcal{I}$ the initial state, and G the goal state. Each action a has sets of preconditions $pre(a)$, add effects $add(a)$, and delete effects $del(a)$. The planning task is to find a sequence of applicable actions that transforms $\mathcal{I}$ into a state containing G.

Figure 2: The sequential plan for the running example (C: Customer; S: Supplier; W: Warehouse).

Given a SCP=(W, C, r_{in}, r_{out}), we translate it into the following classical planning problem.

Actions. We develop actions in A in two steps. For each action $a \in A$, we also define two properties: $key(a)$ denotes the operation or contingency a is representing; and $host(a)$ denotes the Web service w that $key(a)$ belongs to.

1) For each Web service $w \in W$, for each operation $op \in w$, we introduce an **operation action** a with $pre(a) = op.I$ and $add(a) = op.O$. We also define $key(a) = op$ and $host(a) = w$.

2) For each contingency $c \in C$, let $c = (op, I^i, pre^i)$, we introduce a **contingency action** b_c and define $pre(b_c) = \{I_i\}$, $add(b_c) = \{cont_c\}$, $key(b_c) = op$, and $host(b_c) = w$ where $op \in w$. Here, $cont_c$ is a special fact introduced for each contingency c. We also modify the action a for op by $pre(a) = pre(a) \setminus \{op.I_i\} \cup \{cont_c\}$. This way, we compile away the contingency on action outcomes.

Facts. The facts include all of the input and output parameters ($op.I$ and $op.O$) for every operation op in all services $w \in W$, and all of the special contingency fact $cont_c$ for $c \in C$.

Initial and goal states. We set $\mathcal{I} = r_{in}$ and $G = r_{out}$.

Example 2. The SCP shown in Example 1 is translated into a planning problem π with $A = \{ReceiveRFQ, ConfirmPO, \cdots, b_{c1}, b_{c2}\}$.

For action a from operation $ReceiveRFQ$, $pre(a) = \{pid, pid\_name\}$, $add(a) = \{pid\_price\}$. The action b_{c1} for contingency $c_1 = (ConfirmPO, po\_avail, \{avail\})$ has $pre(b_{c1}) = \{po\_avail\}$ and $add(b_{c1}) = \{cont_{c1}\}$. The action a' for $ConfirmPO$ has $pre(a') = \{cont_{c1}\}$ and $add(a') = \{po\_accept, shiporderR\}$.

$F = \{pid, pid\_name, po\_avail, \cdots, cont_{c1}, cont_{c2}\}$, $\mathcal{I} = \{pid, pid\_name, pid\_quantity\}$, and finally, $G = \{purchase\_confirm, po\_order\_reject\}$. ∎

After compilation, we use a planner to find a solution plan for π, which is a sequence of actions $A = (a_1, \cdots, a_m)$ transforming $\mathcal{I}$ to G. A solution to π in Example 2 found by FF (Hoffmann and Nebel 2001) is in Figure 2.

Generation of the master plan

Given a solution plan $A = (a_1, \cdots, a_m)$, we build a dependency graph as follows.

Definition 8 (Dependency). *Given a solution plan $A = (a_1, \cdots, a_m)$, an action a_j depends on a_i (denoted as $a_i \vdash a_j$) if and only if, $i < j$ and there exists a fact $f \in pre(a_j)$ such that $f \notin \mathcal{I}$ and a_i is the last action in $a_1, \cdots, a_{j-1}$ such that $f \in add(a_i)$.*
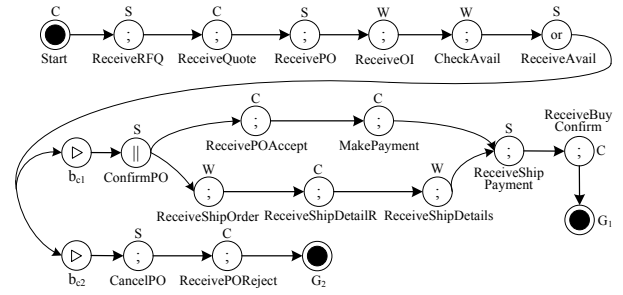

Figure 3: The DG for the example. The goal states (G_1 and G_2) are also added to the graph for better illustration.

Note that the above definition is general for both operation actions and contingency actions.

Definition 9 (Dependency Graph (DG)). *Given a solution plan $A = (a_1, \cdots, a_m)$, a dependency graph is a directed graph $DG = (V, E)$ such that $V = A$ and there is an edge $(a_i, a_j) \in E$ if and only if $a_i \vdash a_j$.*

Example 3. Figure 3 shows the DG for the plan in Figure 2. The DG discovers structures, including sequential, parallel and conditional ones, of the master plan. We observe that there is a choice of making a decision based on contingency actions (b_{c1} and b_{c2}) after the Supplier receives the availability information of an order. Moreover, once a product order is accepted ($ConfirmPO$), roles work in parallel in order to make payment and arrange shipping. ∎

A choreography master plan can be derived from the DG.

- For every operation action a in the DG whose out degree is 1, we mark a by ";".

- For every contingency action b_c in the DG, we mark it by the contingency sign "▷".

- For every action a whose out degree is more than 1, we mark a by "‖" if its successors do not include contingency action and mark a by "or" otherwise.

- For every edge (a_i, a_j), if $host(a_i) \neq host(a_j)$, we mark the edge by $talk(i, j)$.

After such marking, we can write out the choreography master plan by viewing the marked DG as the parsing graph for the master plan language.

Example 4. Figure 3 shows the marking of DG (*talk*'s are not shown). The master plan derived from the DG is:
ReceiveRFQ; *talk*(S,C); ReceiveQuote; $\cdots$; ($c_1 \triangleright$ (ConfirmPO; (*talk*(S,C); ReceivePOAccept; $\cdots$) ‖ (*talk*(S,W); ReceiveShipOrder; $\cdots$); ReceiveShipPayment; *talk*(S, C); ReceiveBuyConfirm)) or ($c_2 \triangleright$ (CancelPO; *talk*(S, C); ReceivePOReject)) ∎

Generation of the distributed plan

To generate the distributed plan, we partition $DG = (V, E)$. We partition the vertex set V into multiple, disjoint sets, one for each role $w \in W$. That is, $V = V_1 \cup \cdots \cup V_N$, where $a \in V_i$ if and only if $host(a) = w_i$. It is illustrated in Figure 4.

Since $a_i \vdash a_j$ only when $i < j$, the DG is acyclic. Hence, a node a_j is an **offspring** of a_i if there is a path from a_i to a_j in DG. We define that an action a_i **depends on** a contingency c if a_i is an offspring of b_c in DG.
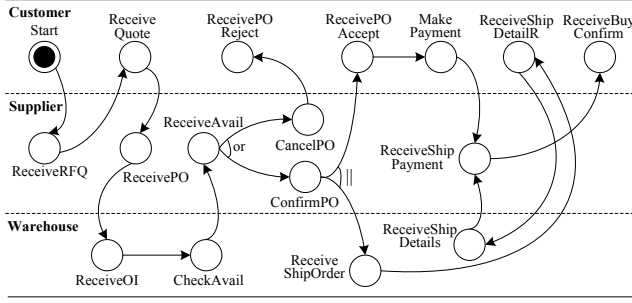
Figure 4: The collaborative interactions between three roles.

Definition 10 (Lead Operation) *Given a solution sequence $A = (a_1, \cdots, a_m)$, for a given contingency $c \in C$ and role $w \in W$, the lead operation $lead(c, w)$ is the first action a in A such that $host(a) = w$ and a depends on c. w may not contain a lead operation if no action in w depends on c.*

Based on the marked DG, for every role we generate a local plan in the language R defined in Section 3. The entire algorithm is long and omitted, and we give the key rules for generation below. For each role $w_k \in W$, $k = 1, \cdots, N$, we consider the actions in partition V_k following the order in A.

- For every two actions a_i and a_j in V_k where $i < j$, if a_j is an offspring of a_i, they are arranged sequentially (";"); if not, they are arranged in parallel ("$\|$") if they depend on the same contingency or by choice ("or") otherwise.

- For every edge from an operation action in V_k to an action in a different partition V_l, we insert $send(ch, k, l)$.

- For every edge from an operation action in a different partition V_l to an action in V_k, we insert $recv(ch, l, k)$.

- For every contingency action b_c in V_k, we insert $send(ch, k, l)$ to every other partition V_l that contains a lead operation.

- For every action a such that $a = lead(c, w_k)$, where $host(b_c) = l$, $l \neq k$, we insert $recv(ch, l, k)$ before a.

Example 5. Based on the partitioned DG in Figure 4, the local plans are, for Customer: $recv$(ch0,S,C); ReceiveQuote; $send$(ch1,C,S); $recv$(ch2,S,C); ($c_2 \triangleright recv$(ch3,S,C); ReceivePOReject) or ($c_1 \triangleright$ (($recv$(ch4,S,C); ReceivePOAccept; MakePayment; $send$(ch5,C,S)) $\|$ ($recv$(ch6,W,C); ReceiveShipDetailR; $send$(ch7,C,W))); $recv$(ch8,S,C); ReceiveBuyConfirm).

Supplier: ReceiveRFQ; $send$(ch0,S,C); $recv$(ch1,C,S); ReceivePO; $send$(ch10,S,W); $recv$(ch11,W,S); ReceiveAvail; ($send$(ch2,S,C) $\|$ $send$(ch8,S,W)); ($c_2 \triangleright$ CancelPO; $send$(ch3,S,C)) or ($c_1 \triangleright$ ConfirmPO; ($send$(ch4,S,C) $\|$ $send$(ch12,S,W)); ($recv$(ch5,C,S) $\|$ $recv$(ch9,W,S)); ReceiveShipPayment; $send$(ch8,S,C)).

Warehouse: $recv$(ch10,S,W); ReceiveQI; CheckAvail; $send$(ch11,W,S); $recv$(ch8,S,W); ($c_1 \triangleright$ ($recv$(ch12,S,W); ReceiveOrder; $send$(ch6,W,C); $recv$(ch7,C,W); ReceiveShipDetails; $send$(ch9,W,S))). ∎

The correctness of our distributed plan can be established by showing its equivalence to the master plan, which satis-

```
1  %customer {
2  method main() =
3    waitfor
4      (request($pid, $pid_name, $pid_quantity) <= peer($enduser)
5      then request($pid, $pid_name) => peer($supplier)
6      then waitfor
7        (reply($pid_price) <= peer($supplier))
8      timeout(call main())
9      then request($pid, $pid_quantity) => peer($supplier)
10     then call wait_decision($pid, $pid_quantity)
11     then call main())
12   timeout(e)
13
14 method wait_decision($pid, $pid_quantity) =
15   waitfor
16     if (reply($po_accept) <= peer($supplier))
17       then (($payment_info = MakePayment($pid, $pid_quanitity)
18       then reply($payment_info) => peer($supplier))
19       par waitfor
20         (request($shiporderR) <= peer($warehouse)
21         then $ship_details = ReceiveShipDeatilR($shiporderR)
22         then reply($ship_details) => peer($warehouse))
23       timeout(e))
24       then waitfor
25         (reply($purchase_confirm) <= peer($supplier))
26       timeout(e)
27     or else (reply($po_order_reject) <= peer($supplier))
28   timeout(e) }
```

Figure 5: Local plan of Customer in MAP specification.

fies the user's need as it solves the planning problem π modeling the logical constraints and contingencies of the SCP.

Theorem 1. *For a SCP, the distributed plan generated by our algorithm is equivalent to the master plan.*

Due to space limit, we only sketch the key idea of proof here. Our proof constructs an one-to-one correspondence between any possible sequence in the master plan to a sequence in the distributed plan, ignoring the communication primitives. (Qiu et al. 2007) studied necessary conditions for such equivalence and concluded that a natural projection from a master plan does not always give an equivalent distributed plan. There is a key difference: their work considers any possible master plan, while in our approach, only those master plans and local plans that can be generated from our planning method are considered. Our planning work, although more restrictive, ensures the equivalence of local and master plans by analyzing the dependency and contingency and using communication to enforce actions' partial orders.

A counterexample in (Qiu et al. 2007) is a master plan $(a_1^1 \| a_2^1); a_1^2$, where a_1^1 and a_2^1 are in role 1 and a_1^2 in role 2. The distributed plan is $a_1^1; a_2^1$ for role 1 and a_1^2 for role 2, which allows $a_1^1; a_2^1; a_1^2$, a sequence the master plan does not allow. In our planning approach, however, such a nonequivalence will not occur. If a_2^1 depends on a_1^2, then there will be a send/recv pair that restricts the order; if a_2^1 does not depend on a_1^2, then the master plan will be $(a_1^1; a_2^1) \| a_1^2$, which is equivalent to the distributed plan.

In addition, under contingencies, it is suggested that equivalence relies on the existence a *dominant role* (Qiu et al. 2007), which makes choices that all other roles will follow. In our method, we essentially have a dominant role for each contingency (whoever generates the output parameter), and the dominant role sends the decision to the lead operations in other roles.

Experimental Results

We developed a prototype system in Java. It takes service repositories in WSDL as input, allows user to specify initial state and choreography goals, and applies our approach to

generate a distributed choreography plan in the MAP specification (Barker, Walton, and Robertson 2009). Two planners FF (Hoffmann and Nebel 2001) and SatPlan06 (Kautz, Selman, and Hoffmann 2006) are integrated in our system. We ran our experiments on a PC with Intel Pentium(R) dual core processor 2.4 GHz and 1G RAM.

Currently a user needs to specify contingencies the first time a Web service is published. In fact, this involves only a slight enhancement to a Web service description language such as WSDL or OWL-S. In our experiment, we have enhanced WSDL by some special annotations to describe the contingencies. Since typically only a small number of Web services involve contingencies, it does not require much work as most Web services do not need any change. Then, the translation from the WSDL repository to the planning formulation is completely automated and does not involve any manual work. That is, our planning translation algorithm will automatically generate the correct planning domain specification in the PDDL language. It can parse the special annotations for contingencies and correctly translate them into contingency actions.

Taking the choreography dependency graph in Figure 3 as an example, we convert it into three MAP specifications. Figure 5 shows the MAP plan for Customer. Supplier and Warehouse are omitted. The translation from our language R to MAP is direct from the example. Details of MAP can be found in (Barker, Walton, and Robertson 2009).

We tested and compared our approach with WSPR (Oh, Lee, and Kumara 2007) on 81,464 Web services in 18 groups of large scale service repositories from the ICEBE05 Web services composition challenge. We compared against WSPR since it is a well-known solver for automated composition of Web services using AI planning techniques, i.e., forward search and regression search algorithm. Note that WSPR solves orchestration, a much simpler problem than the choreography problem that our approach solves. Since the most time-consuming components in our approach are SCP translation and choreography plan generation, we report the timing of these two parts in Table 1. The response time of finding a choreography plan is the duration from submitting a chorography task until receiving the output of the final solution or failing to find a plan.

Although the most difficult repository with the largest number of services and I/O parameters (Composition2-100-32) takes more than 1,100 seconds for the SCP translation, it can be performed offline only once. Each time when a user wants to find a service choreography, he only needs to specify the inputs and possible goals. The system will automatically generate the adjusted planning formulation very quickly. In other words, once we translate operations in a large Web service repository into planning actions, we do not need to parse a service repository again when a user submits a choreography task.

Comparing to WSPR, we can see that although our approach solves the more complex choreography problem and handles distributed coordination, communication, and contingency, it is more efficient than WSPR in finding a solution plan on the average response time. The reason is that our translation to PDDL allows us to leverage on the advances

Table 1: The SCP translation time and average response time of finding a choreography plan using FF and SatPlan06 and WSPR. There are 18 groups of service repositories on the ICEBE05. Each repository has 11 choreography goals. Column '|P|' is the size of input and output parameters in an operation. Column '|W|' is the number of services. Column 'Trans. time' gives the time for SCP translation. Column 'Generation time' is the average response time of finding a plan for all 11 choreography goals in seconds.

Repository	|P|	|W|	Trans. time	Generation time		
				FF	SatPlan06	WSPR
Composition1-20-4		2156	74.03	0.306	0.800	8.674
Composition1-50-4	4-8	2656	109.77	0.344	0.980	11.242
Composition1-100-4		4156	250.92	0.491	1.539	17.665
Composition1-20-16		2156	75.641	1.149	2.334	17.753
Composition1-50-16	16-20	2656	112.25	1.198	3.565	22.478
Composition1-100-16		4156	253.00	1.873	5.419	36.278
Composition1-20-32		2156	77.20	3.224	5.076	29.988
Composition1-50-32	32-36	2656	112.95	3.422	7.910	37.726
Composition1-100-32		4156	260.172	4.792	14.221	62.629
Composition2-20-4		3356	164.92	0.569	2.672	14.878
Composition2-50-4	4-8	5356	400.14	0.845	4.391	24.046
Composition2-100-4		8356	964.03	1.214	7.373	38.934
Composition2-20-16		6712	625.45	4.359	30.210	63.430
Composition2-50-16	16-20	5356	412.31	3.817	18.307	49.776
Composition2-100-16		8356	972.47	5.320	32.235	81.791
Composition2-20-32		3356	175.74	7.819	26.609	50.719
Composition2-50-32	32-36	5356	417.44	11.098	49.489	86.794
Composition2-100-32		8356	1,103.38	14.833	92.536	148.807

of automated planners, while WSPR uses its own planning model and composition solver.

Conclusions

Automatic and efficient Web service composition (WSC) can potentially simplify the implementation of business processes. Web service choreography is an important paradigm for WSC with many advantages over Web service orchestration. This paper presents a method to generate distributed plans for Web service choreography. Based on an AI planning approach and dependency graph analysis, the method is efficient and fully automated, and also ensures correct collaboration of multiple peers.

In this paper, we focus on finding a feasible choreography plan that can correctly achieve choreography goals under distributed collaboration and contingencies. For our future work, we will further consider the optimization of choreography plan quality, by considering important non-functional property metrics, such as multiple QoS of Web services.

Acknowledgment

References

Agarwal, V.; Chafle, G.; Dasgupta, K.; et al. 2005. Synthy: A system for end to end composition of Web services. *JWS* 3(4):311–339.

Barker, A.; Walton, C. D.; and Robertson, D. 2009. Choreographing Web services. *IEEE Trans. Serv. Comput.* 2(2):152–166.

Hoffmann, J., and Nebel, B. 2001. The FF planning system: fast plan generation through heuristic search. *J. of AI Research* 14(1):253–302.

Hoffmann, J.; Bertoli, P.; and Pistore, M. 2007. Web service composition as planning, revisited: In between background theories and initial state uncertainty. In *Proc. of AAAI*.

Kang, Z.; Wang, H.; and Hung, P. 2007. WS-CDL+: an extended WS-CDL execution engine for Web service collaboration. In *Proc. of ICWS*.

Kautz, H.; Selman, B.; and Hoffmann, J. 2006. SatPlan: Planning as satisfiability. In *Abstracts IPC5*.

Oh, S. C.; Lee, D.; and Kumara, S. R. T. 2007. Web Service Planner(WSPR): An effective and scalable Web service composition algorithm. *Int. J. Web Serv. Res.* 4(1):1–22.

Peltz, C. 2003. Web services orchestration and choreography. *Computer* 36(10):46–52.

Ponnekanti, S. R., and Fox, A. 2002. Sword: A developer toolkit for web service composition. In *Proc. of WWW*.

Qiu, Z.; Zhao, X.; Cai, C.; and Yang, H. 2007. Towards the theoretical foundation of choreography. In *Proc. of WWW*.

Yang, H.; Zhao, X.; Cai, C.; and Qiu, Z. 2008. Model-checking of web services choreography. In *Proc. of SOSE*.